

[Github.com/KnaveAndVarlet/ADASS2024\\_GPU](https://github.com/KnaveAndVarlet/ADASS2024_GPU)

KnaveAndVarlet / ADASS2024\_GPU

Type / to search

Code Issues Pull requests Actions Projects Wiki Security Insights Settings

ADASS2024\_GPU Public

Pin

Unwatch 1

Fork 0

Star 1

main

1 Branch

4 Tags

Go to file

Add file

Code

About



KnaveAndVarlet Update README.md

b79d980 · 2 days ago 10 Commits

Metal

Fixes issues found building on some Linux systems.

2 days ago

Vulkan

Fixes issues found building on some Linux systems.

2 days ago

GPU\_Tutorial\_Instructions.pdf

Fixes issues found building on some Linux systems.

2 days ago

GPU\_Tutorial\_Instructions\_Windows.pdf

Added Makefiles and Instructions for Windows.

last week

LICENSE

Update LICENSE

last month

README.md

Update README.md

2 days ago

README License

# ADASS2024\_GPU

Example programs for the ADASS 2024 tutorial on programming laptop GPUs.

Note: The tutorial instructions on the ADASS website may not always reflect the very latest changes shown in the versions here.

Pin

Unwatch 1

Fork 0

Star 1

About



Example programs for the ADASS 2024 tutorial on programming laptop GPUs.

Readme

View license

Activity

1 star

1 watching

0 forks

Releases 4

v1.1.1 Minor bug fixes Latest

2 days ago

+ 3 releases

Packages

No packages published

[Publish your first package](#)

Languages



C++ 98.0%

Python 0.5%

```
sudo apt update  
(sudo apt install build-essential)  
sudo apt install vulkan-tools  
sudo apt install libvulkan-dev  
sudo apt install vulkan-validationlayers-dev spirv-tools  
sudo apt install libglfw3-dev
```

Ubuntu

Minimal Vulkan and GLFW installation  
(and C++ tools)

```
sudo dnf check-update  
sudo dnf upgrade  
(sudo dnf install gcc-c++)  
(sudo dnf install make)  
sudo dnf install vulkan-tools  
sudo dnf install vulkan-loader-devel  
sudo dnf install mesa-vulkan-devel vulkan-validation-layers-devel  
sudo dnf install glfw-devel
```

Fedora

```
float gridXcent = gridSize.x * 0.5;
float gridYcent = gridSize.y * 0.5;
float x0 = args->xCent + (index2.x - gridXcent) * args->dX;
float y0 = args->yCent + (index2.y - gridYcent) * args->dY;

float x = 0.0;
float y = 0.0;
uint iteration = 0;
uint max_iteration = args->iter;
float xtmp = 0.0;
while(((x * x) + (y * y) <= 4.0) && (iteration < max_iteration))
{
    xtmp = (x + y) * (x - y) + x0;
    y = (2.0 * x * y) + y0;
    x = xtmp;
    iteration += 1;
}

float value = iteration;
if (iteration >= max_iteration) value = 0.0;
uint index = index2.y * gridSize.x + index2.x;
out[index] = value;
```

Metal code

```
float gridXcent = args.nx * 0.5;
float gridYcent = args.ny * 0.5;
float x = args.xCent + (float(gl_GlobalInvocationID.x) - gridXcent) * args.dX;
float y = args.yCent + (float(gl_GlobalInvocationID.y) - gridYcent) * args.dY;

vec2 c = vec2(x,y);
vec2 z = vec2(0.0,0.0);
float n = 0.0;
const int M = args.iter;
for (int i = 0; i<M; i++)
{
    n++;
    z = vec2(z.x*z.x - z.y*z.y, 2.*z.x*z.y) + c;
    if (dot(z, z) > 4.0) break;
}
if (n >= M) n = 0.0;

imageData[args.nx * gl_GlobalInvocationID.y + gl_GlobalInvocationID.x] = n;
```

GLSL code

```
float InArray[NY][NX];
float OutArray[NY][NX];

for (int Iy = 0; Iy < NY; Iy++) {
    for (int Ix = 0; Ix < NX; Ix++) {
        OutArray[Iy][Ix] = InArray[Iy][Ix] + Ix + Iy;
    }
}
```

```
float InArray[NY][NX];
float OutArray[NY][NX];

for (int Iy = 0; Iy < NY; Iy++) {
    for (int Ix = 0; Ix < NX; Ix++) {
        OutArray[Iy][Ix] = InArray[Iy][Ix] + Ix + Iy;
    }
}
```

```
float* In = new float[Nx * Ny];
float* Out = new float[Nx * Ny];

for (int Iy = 0; Iy < Ny; Iy++) {
    for (int Ix = 0; Ix < Nx; Ix++) {
        Out[Iy * Nx + Ix] = In[Iy * Nx + Ix] + Ix + Iy;
    }
}
```

```
kernel void adder(device float *in [[ buffer(1) ]],  
                  device float *out[[ buffer(2) ]],  
                  uint2 index2 [[thread_position_in_grid]],  
                  uint2 gridSize [[threads_per_grid]]) {  
  
    uint nx = gridSize.x;  
    uint ix = index2.x;  
    uint iy = index2.y;  
    out[iy * nx + ix] = in[iy * nx + ix] + iy + ix;  
}
```

Metal Adder GPU code

```
struct AdderArgs {int nx; int ny; };
layout(std140,binding = 0) uniform paramBuf { AdderArgs args; };
layout(binding = 1) readonly buffer inBuf { float inputData[]; };
layout(binding = 2) writeonly buffer outBuf { float outputData[]; };

void main() {

    uint ix = gl_GlobalInvocationID.x;
    uint iy = gl_GlobalInvocationID.y;
    uint nx = args.nx;
    uint ny = args.ny;
    if (ix >= nx || iy >= ny) return;
    outputData[iy * nx + ix] = inputData[iy * nx + ix] + ix + iy;
}
```

Vulkan Adder GPU code

# GRAPHICS REINVENTED

NVIDIA's newest flagship graphics card is a revolution in gaming realism and performance. Its powerful NVIDIA Turing™ GPU architecture, breakthrough technologies, and 11 GB of next-gen, ultra-fast GDDR6 memory make it the world's ultimate gaming GPU.

GeForce RTX™ graphics cards are powered by the Turing GPU architecture and the all-new RTX platform. This gives you up to 6x the performance of previous-generation graphics cards and brings the power of real-time ray tracing and AI to games.

When it comes to next-gen gaming, it's all about realism. GeForce RTX 2080 Ti is light years ahead of other cards, delivering truly unique real-time ray-tracing technologies for cutting-edge, hyper-realistic graphics.

## KEY FEATURES

- Real-Time Ray Tracing
- NVIDIA® GeForce Experience™
- NVIDIA® Ansel
- NVIDIA® Highlights
- NVIDIA® G-SYNC™ Compatible



GeForce RTX™ 2080 Ti 11GB

## SPECIFICATIONS

### NVIDIA GeForce RTX™ 2080 Ti 11GB Blower

PNY Part Number	<b>VCG2080T11BLMPB</b>
UPC Code	<b>751492621722</b>
Card Dimensions	<b>1.38" x 10.50" x 4.50", Dual-Slot</b>
Box Dimensions	<b>6.76" x 12.21" x 3.53"</b>
NVIDIA® CUDA Cores	<b>4352</b>
Clock Speed	<b>1350 MHz</b>
Boost Speed	<b>1545 MHz</b>
Memory Speed (Gbps)	<b>14</b>
Memory Size	<b>11GB GDDR6</b>
Memory Interface	<b>352-bit</b>
Memory Bandwidth (GB/sec)	<b>616</b>
TDP	<b>250 W<sup>5</sup></b>
NVLink	<b>2-way</b>
Outputs	<b>DisplayPort 1.4 (x3), HDMI 2.0b, USB Type-C</b>
Multi-Screen	<b>Yes</b>
Resolution	<b>7680 x 4320 @60Hz (Digital)<sup>1</sup></b>
Power Input	<b>Two 8-Pin</b>
Bus Type	<b>PCI-Express 3.0 x16</b>

## Chip



### Apple M2 chip

8-core CPU with  
four performance cores  
and four efficiency cores  
**10-core GPU**  
16-core Neural Engine  
100GB/s memory  
bandwidth



### Apple M2 chip

8-core CPU with  
four performance cores  
and four efficiency cores  
**10-core GPU**  
16-core Neural Engine  
100GB/s memory  
bandwidth



### Apple M2 Pro chip

10-core CPU with  
six performance cores and  
four efficiency cores  
**16-core GPU**  
16-core Neural Engine  
200GB/s memory  
bandwidth

## Media engine

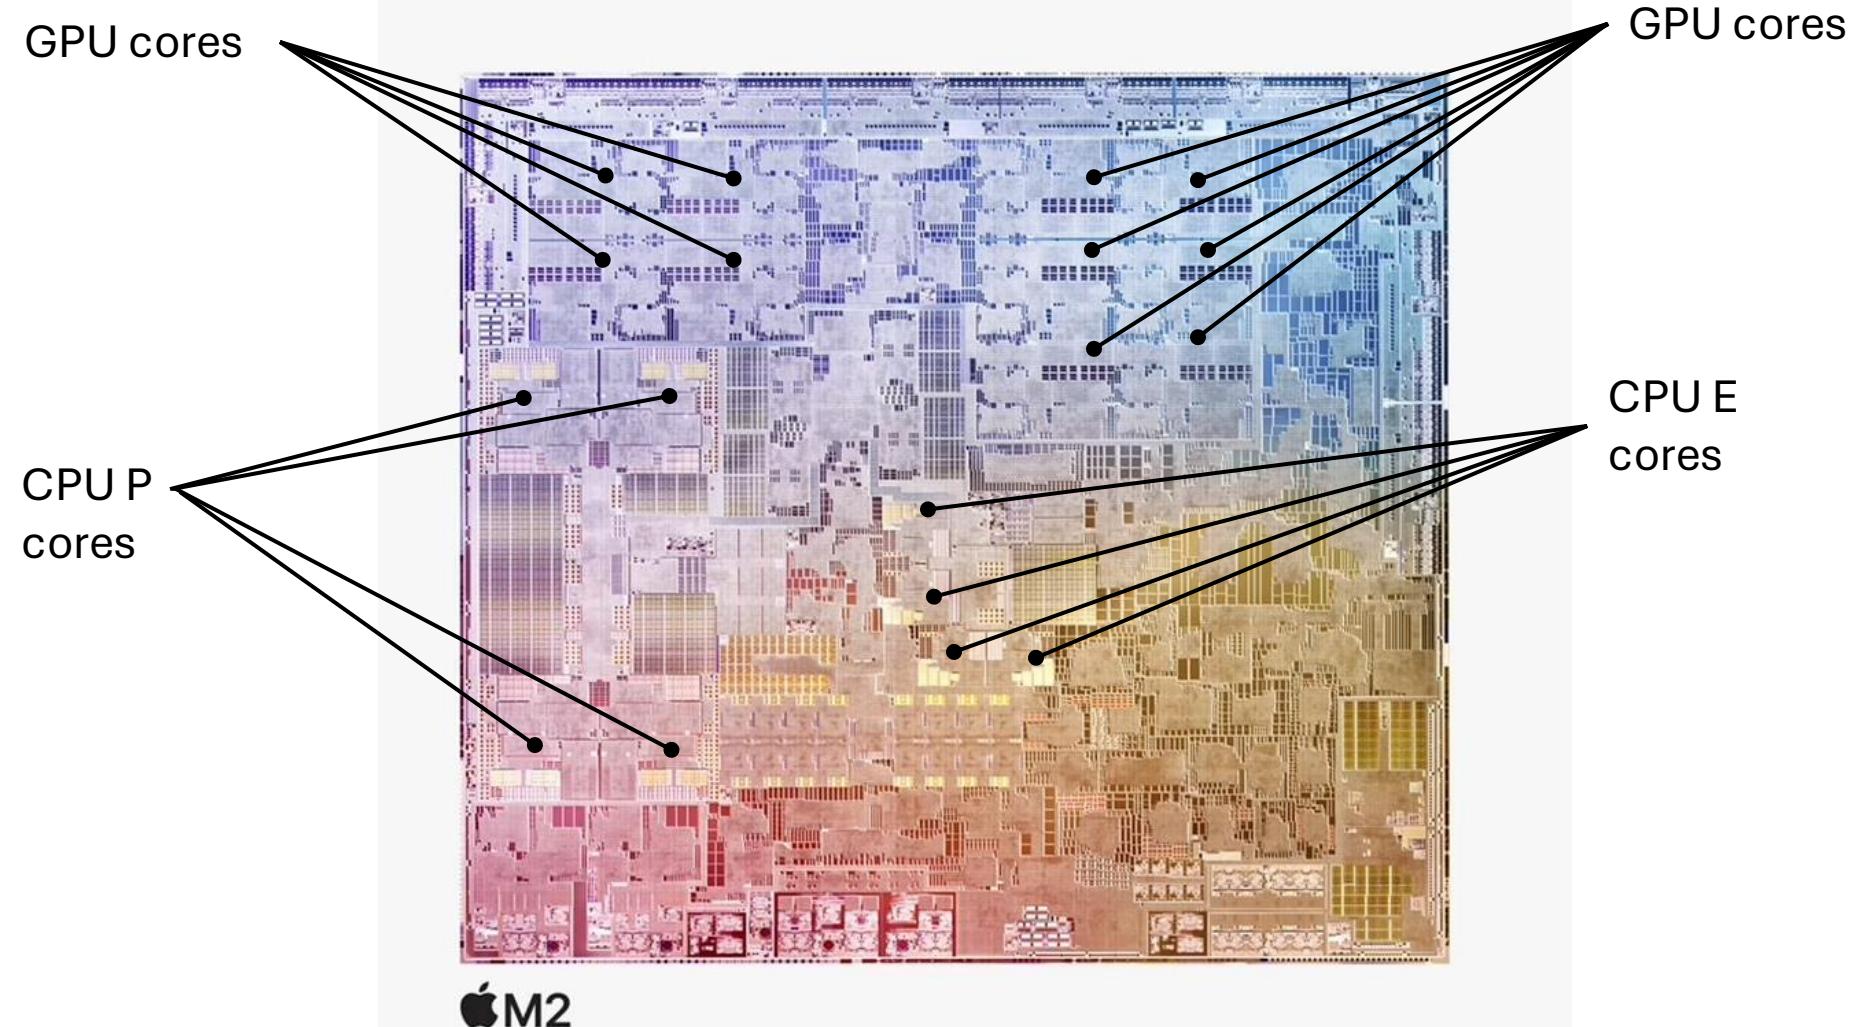
Hardware-accelerated  
H.264, HEVC, ProRes and  
ProRes RAW  
Video decode engine  
Video encode engine  
ProRes encode and  
decode engine

## Media engine

Hardware-accelerated  
H.264, HEVC, ProRes and  
ProRes RAW  
Video decode engine  
Video encode engine  
ProRes encode and  
decode engine

## Media engine

Hardware-accelerated  
H.264, HEVC, ProRes and  
ProRes RAW  
Video decode engine  
Video encode engine  
ProRes encode and  
decode engine

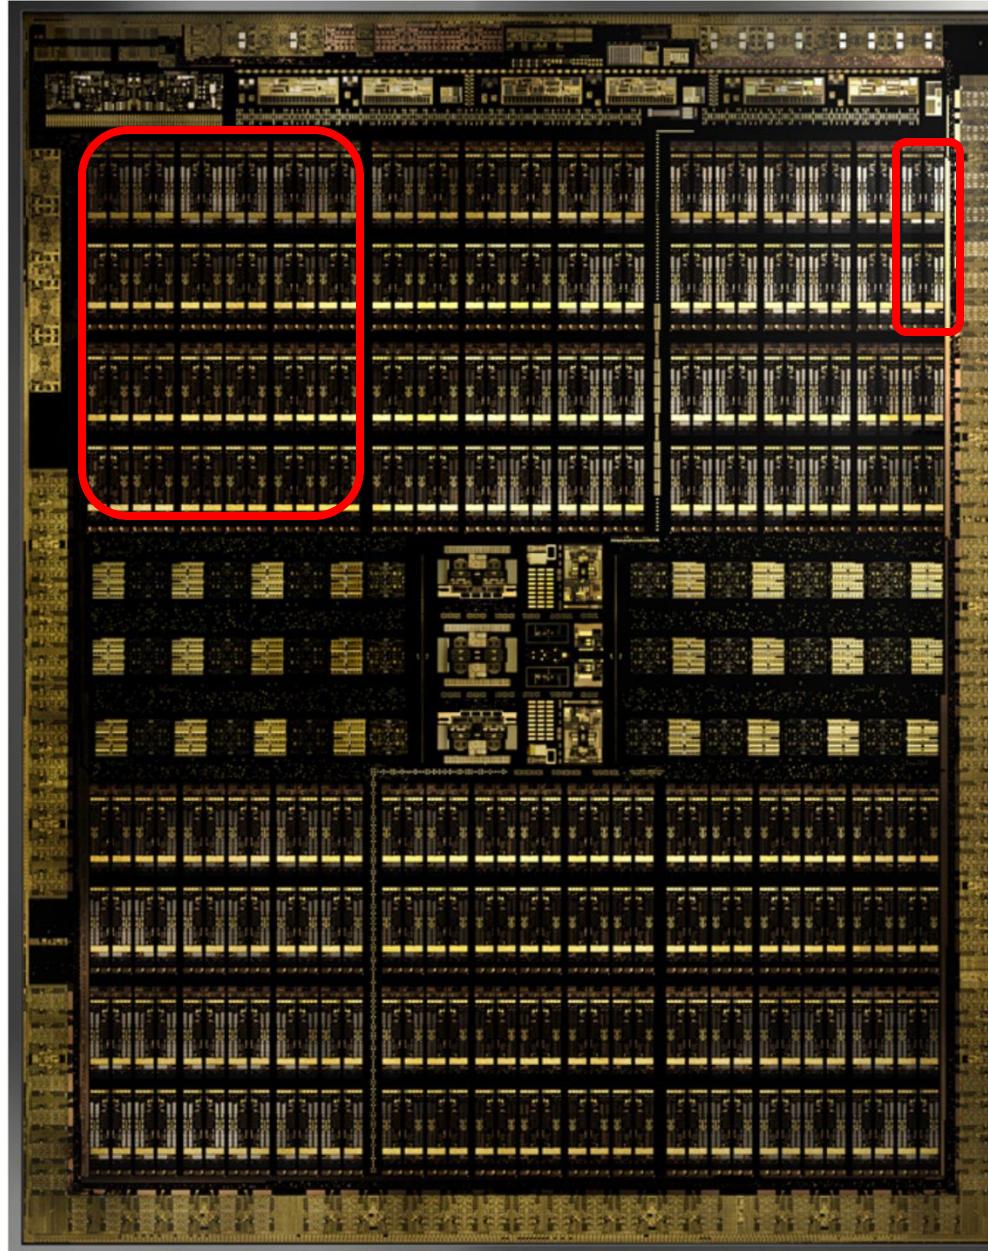


Apple Silicon M2

GPC – GPU Processing Cluster, with 12 SMs (Streaming Multiprocessors) each.

6 GPCs per device.

Total of 72 SMs per device.



Streaming multiprocessor

NVIDIA Turing TU102 GPU



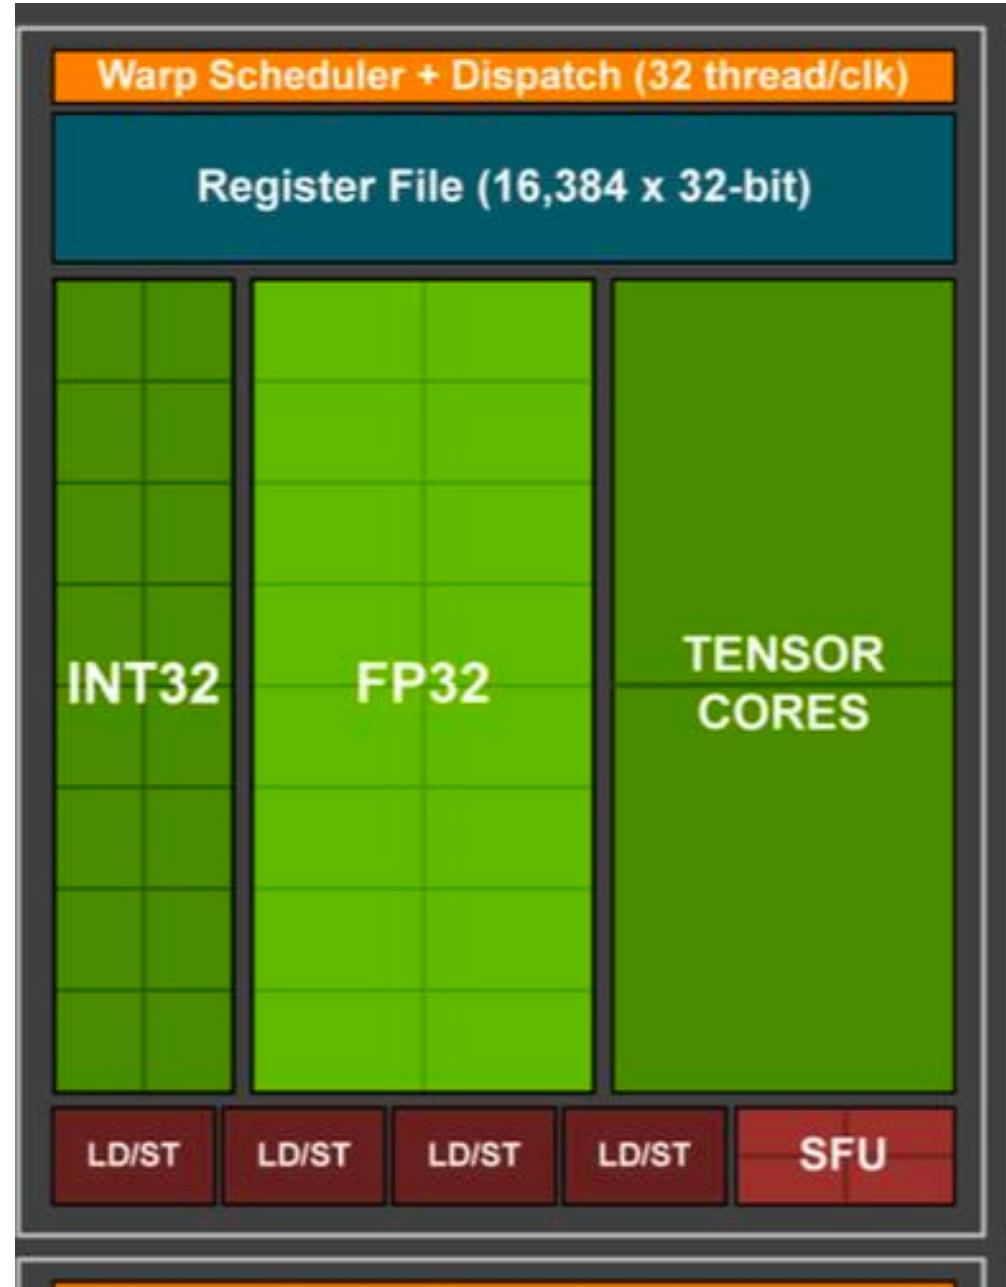
Note: The TU102 GPU also features 144 FP64 units (two per SM), which are not depicted in this diagram. The FP64 TFLOP rate is 1/32nd the TFLOP rate of FP32 operations. The small number of FP64 hardware units are included to ensure any programs with FP64 code operates correctly.

**Figure 2. Turing TU102 Full GPU with 72 SM Units**



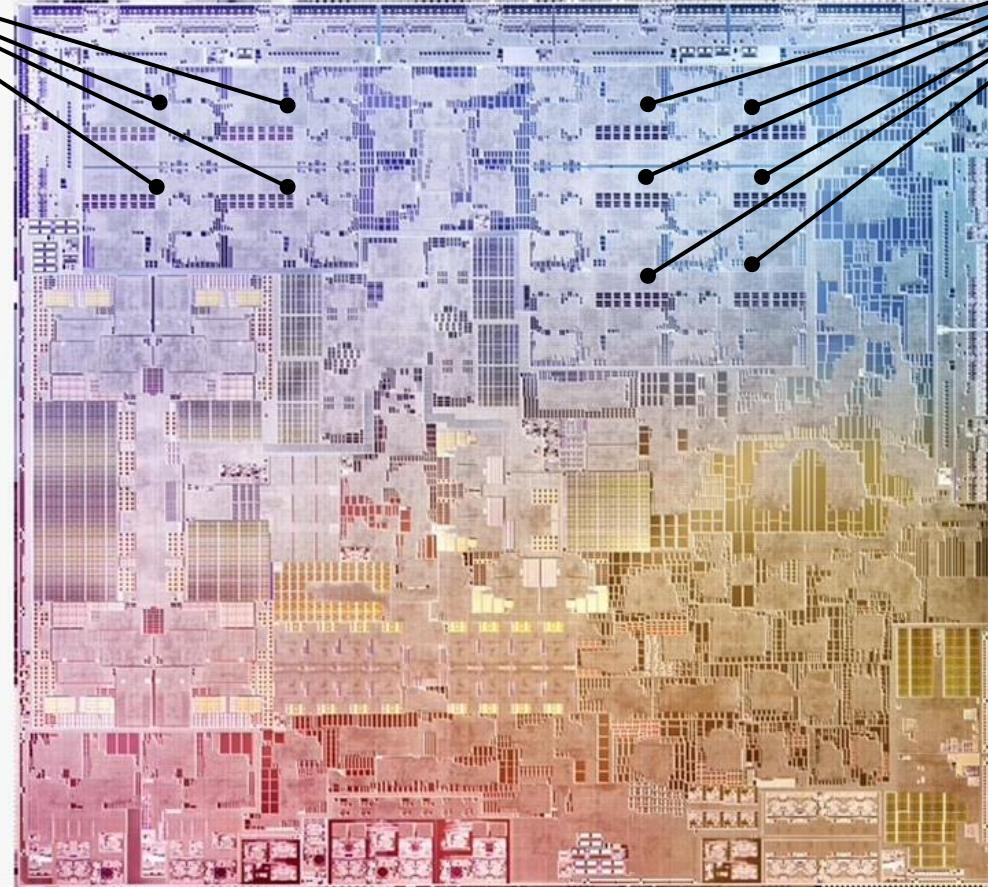
**Figure 4. Turing TU102/TU104/TU106 Streaming Multiprocessor (SM)**

Nvidia TU102  
Streaming multiprocessor (SM)  
sub-partition (SMSP)



GPU cores

GPU cores



Apple M2

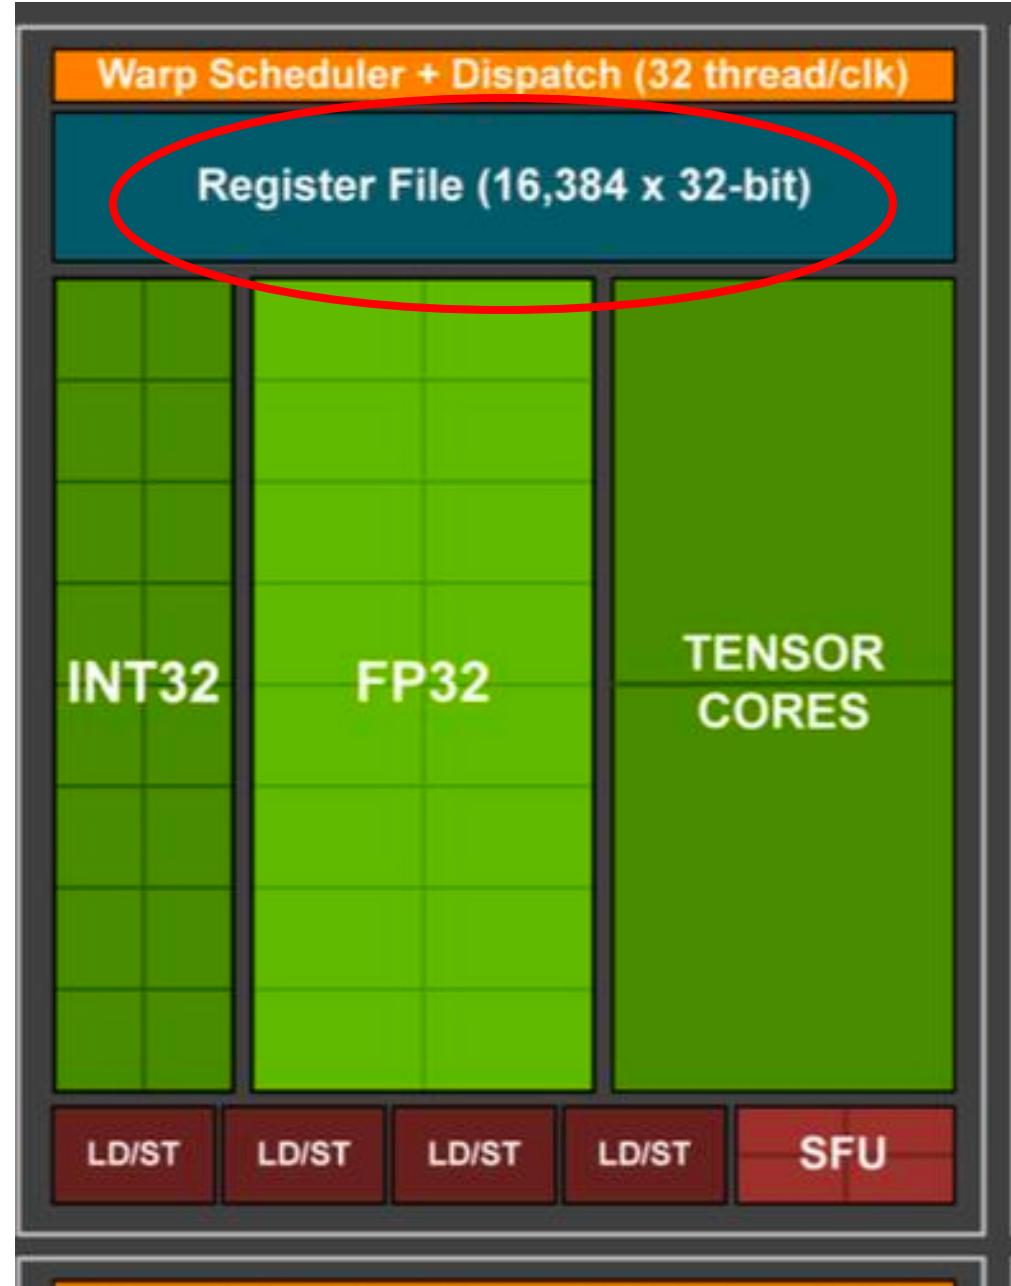
Apple Silicon M2

Each of the 10 GPU  
core has  
16 'execution units'  
each with 8 ALUs.

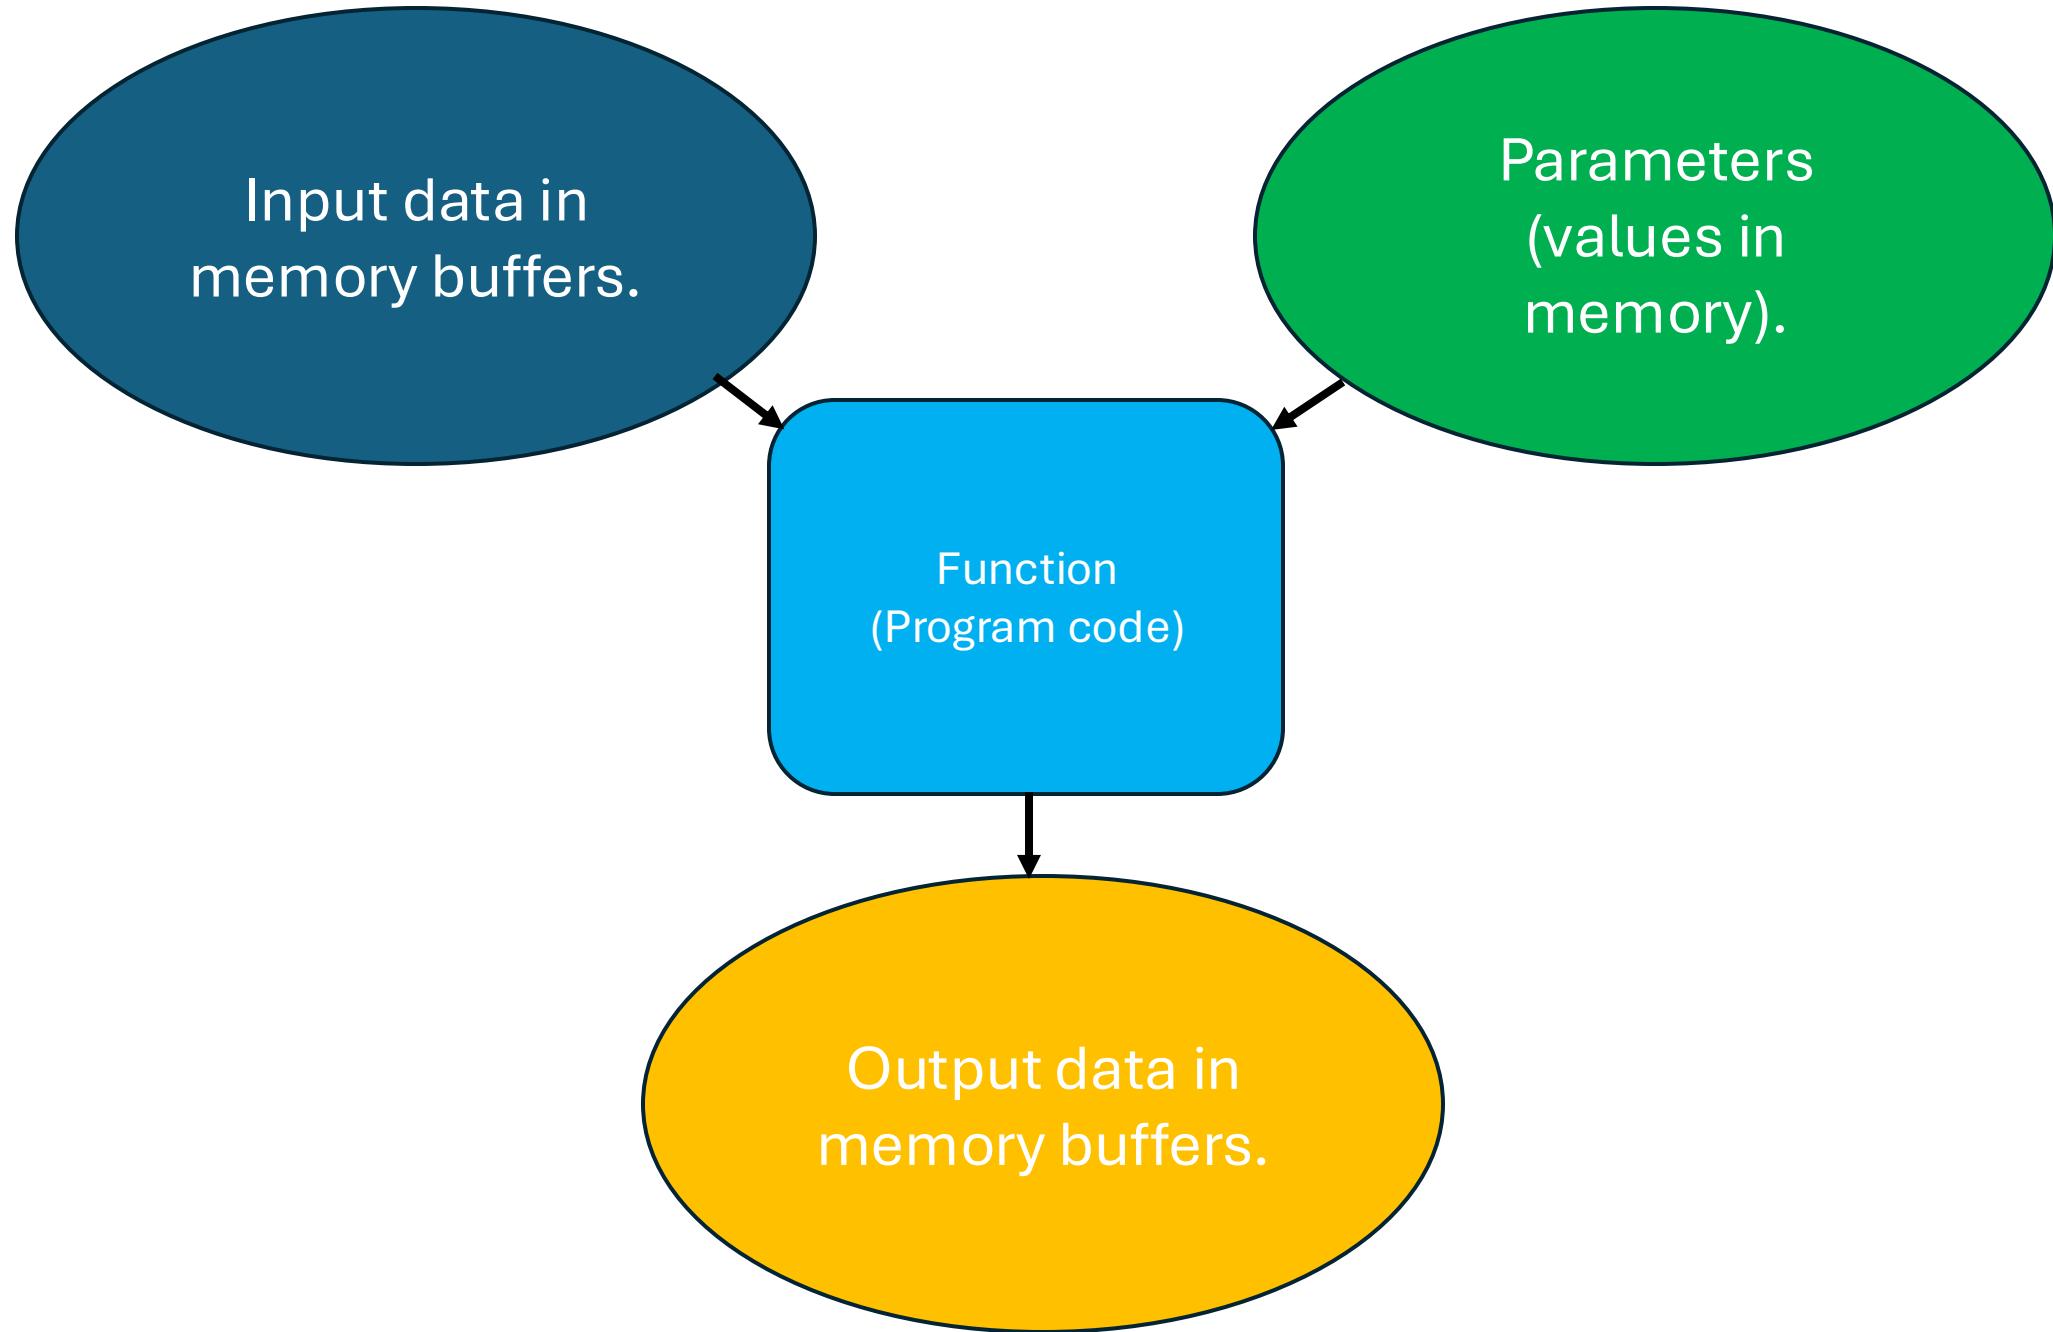
$$10 \times 16 \times 8 = 1280 \text{ ALUs.}$$

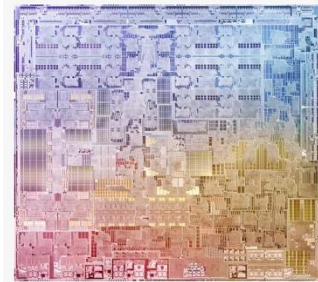
Nvidia TU102  
Streaming multiprocessor (SM)  
sub-partition (SMSP)

Runs threads in groups  
of 32 at a time.  
32 threads is a 'Warp'.



```
MTL::Device* Device = MTLCreateSystemDefaultDevice();
MTL::Library* Library = Device->newLibrary(NS::String::string("Compute.metallib",
                                                               UTF8StringEncoding),&ErrorPtr);
AdderFunction = Library->newFunction(NS::String::string("adder",UTF8StringEncoding));
MTL::Buffer* InputBuffer = Device->newBuffer(AllocationSize,BufferOptions);
MTL::Buffer* OutputBuffer = Device->newBuffer(AllocationSize,BufferOptions);
MTL::CommandQueue* CommandQueue = Device->newCommandQueue();
MTL::ComputePipelineState* PipelineState =
    Device->newComputePipelineState(AdderFunction,&ErrorPtr);
MTL::Size GridSize(Nx,Ny,1);
MTL::Size ThreadGroupDims(ThreadGroupSize / ThreadWidth,ThreadWidth,1);
MTL::CommandBuffer* CommandBuffer = CommandQueue->commandBuffer();
MTL::ComputeCommandEncoder* Encoder = CommandBuffer->computeCommandEncoder();
Encoder->setComputePipelineState(PipelineState);
Encoder->setBuffer(InputBuffer,0,1);
Encoder->setBuffer(OutputBuffer,0,2);
Encoder->dispatchThreads(GridSize,ThreadGroupDims);
Encoder->endEncoding();
CommandBuffer->commit();
CommandBuffer->waitUntilCompleted();
```

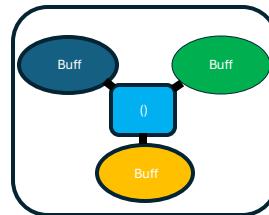




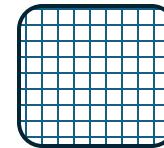
1 : Access GPU device



2 : Create memory buffers



5: Set thread layout



Q

3 : Create GPU program(s)

4 : Create Pipeline(s)

6 : Create command queue(s)

CB

CB

Q CB

7 : Create Command buffer

8 : Encode Command buffer

9: Run command buffer

```
MTL::Device* Device = MTLCreateSystemDefaultDevice();
MTL::Library* Library = Device->newLibrary(NS::String::string("Compute.metallib",
                                                               UTF8StringEncoding),&ErrorPtr);
AdderFunction = Library->newFunction(NS::String::string("adder",UTF8StringEncoding));
MTL::Buffer* InputBuffer = Device->newBuffer(AllocationSize,BufferOptions);
MTL::Buffer* OutputBuffer = Device->newBuffer(AllocationSize,BufferOptions);
MTL::CommandQueue* CommandQueue = Device->newCommandQueue();
MTL::ComputePipelineState* PipelineState =
    Device->newComputePipelineState(AdderFunction,&ErrorPtr);
MTL::Size GridSize(Nx,Ny,1);
MTL::Size ThreadGroupDims(ThreadGroupSize / ThreadWidth,ThreadWidth,1);
MTL::CommandBuffer* CommandBuffer = CommandQueue->commandBuffer();
MTL::ComputeCommandEncoder* Encoder = CommandBuffer->computeCommandEncoder();
Encoder->setComputePipelineState(PipelineState);
Encoder->setBuffer(InputBuffer,0,1);
Encoder->setBuffer(OutputBuffer,0,2);
Encoder->dispatchThreads(GridSize,ThreadGroupDims);
Encoder->endEncoding();
CommandBuffer->commit();
CommandBuffer->waitUntilCompleted();
```

```

KVulkanFramework Framework;
Framework.CreateVulkanInstance(StatusOK);
Framework.FindSuitableDevice(StatusOK);
Framework.CreateLogicalDevice(StatusOK);
InputBufferHndl = Framework.SetBufferDetails(C_InputBufferBinding, "STORAGE", "SHARED", StatusOK);
Framework.CreateBuffer(InputBufferHndl, Length, StatusOK);
float* InputBufferAddr = (float*)Framework.MapBuffer(InputBufferHndl, &Bytes, StatusOK);
OutputBufferHndl = Framework.SetBufferDetails(C_OutputBufferBinding, "STORAGE",
                                              "SHARED", StatusOK);
Framework.CreateBuffer(OutputBufferHndl, Length, StatusOK);
float* OutputBufferAddr = (float*)Framework.MapBuffer(OutputBufferHndl, &Bytes, StatusOK);
UniformBufferHndl = Framework.SetBufferDetails(C_UniformBufferBinding,
                                                "UNIFORM", "SHARED", StatusOK);
Framework.CreateBuffer(UniformBufferHndl, SizeInBytes, StatusOK);
void* UniformBufferAddr = Framework.MapBuffer(UniformBufferHndl, &Bytes, StatusOK);
Framework.CreateVulkanDescriptorSetLayout(Handles, &SetLayout, StatusOK);
Framework.CreateVulkanDescriptorPool(Handles, 1, &DescriptorPool, StatusOK);
Framework.AllocateVulkanDescriptorSet(SetLayout, DescriptorPool, &DescriptorSet, StatusOK);
Framework.SetupVulkanDescriptorSet(Handles, DescriptorSet, StatusOK);
Framework.CreateComputePipeline("Adder.spv", "main",
                               &SetLayout, &ComputePipelineLayout, &ComputePipeline, StatusOK);
Framework.GetDeviceQueue(&ComputeQueue, StatusOK);
Framework.CreateCommandPool(&CommandPool, StatusOK);
Framework.CreateComputeCommandBuffer(CommandPool, &CommandBuffer, StatusOK);
Framework.RecordComputeCommandBuffer(CommandBuffer, ComputePipeline,
                                      ComputePipelineLayout, &DescriptorSet, WorkGroupCounts, StatusOK);
Framework.RunCommandBuffer(ComputeQueue, CommandBuffer, StatusOK);

```

Vulkan ‘Adder’ code - skeleton



Image 1024 x 512 pixels

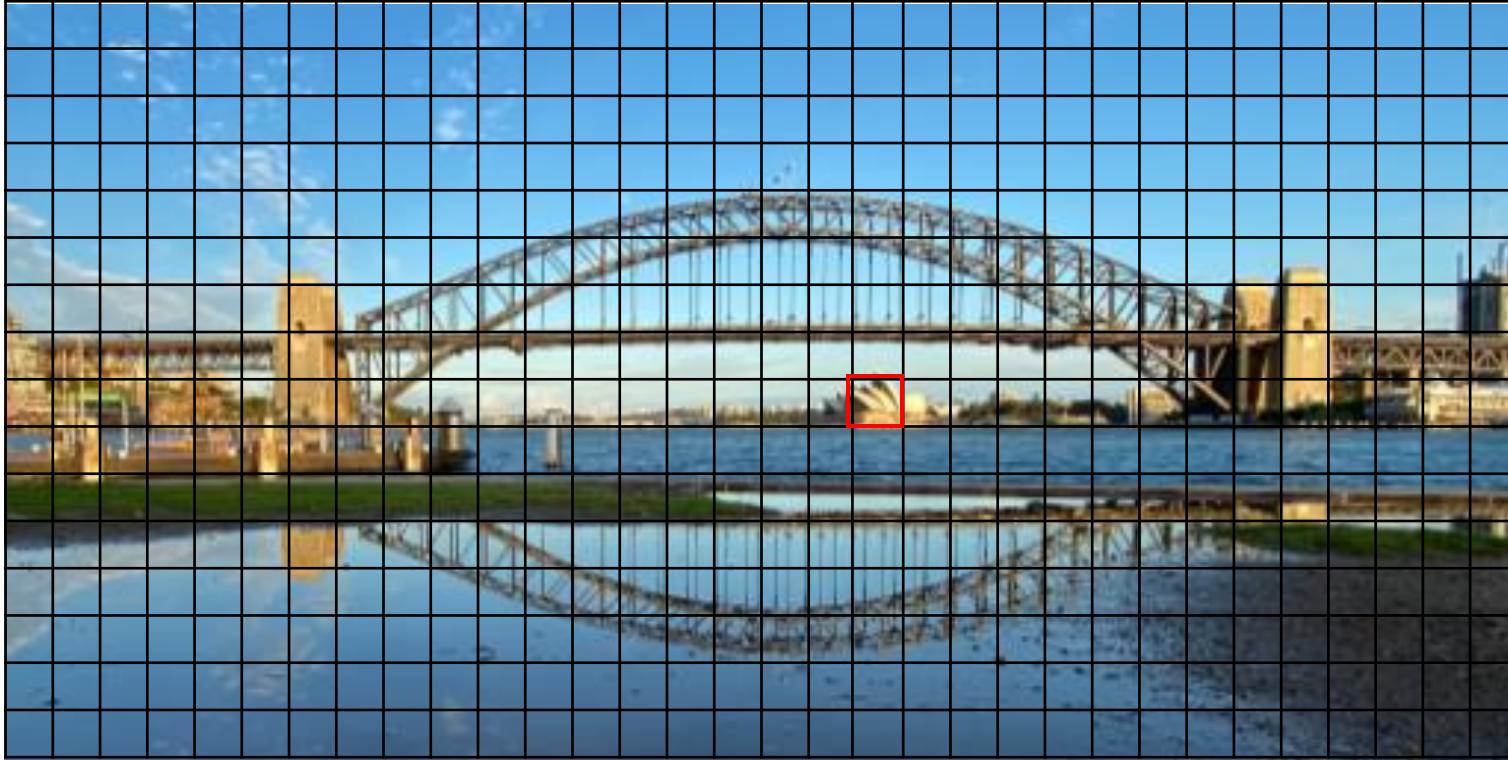
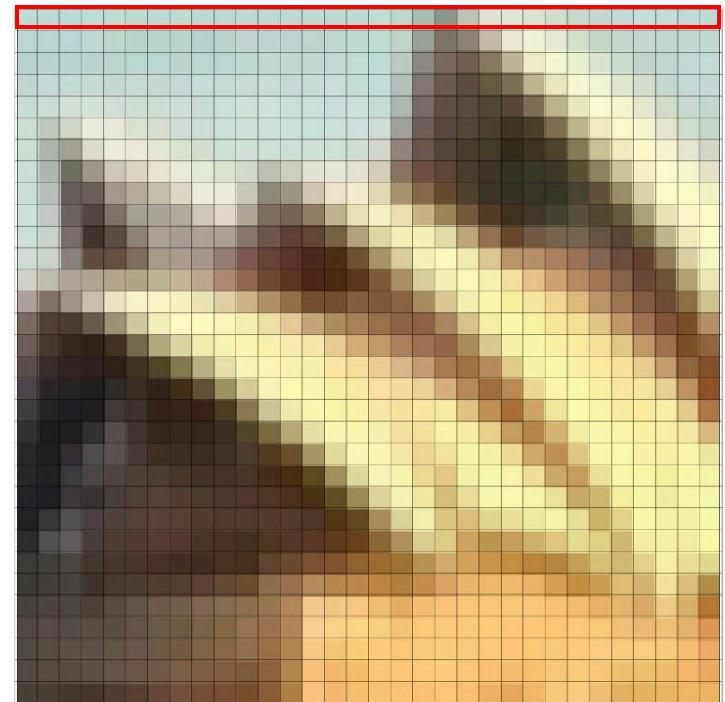


Image 1024 x 512 pixels

Split into 32 x 16 ‘threadgroups’,  
each of 32 x 32 threads, in this case.

Each thread in a threadgroup  
works on one image pixel, in  
this case.

Threads in a threadgroup  
execute in ‘warps’ of 32 threads  
running together in lockstep.



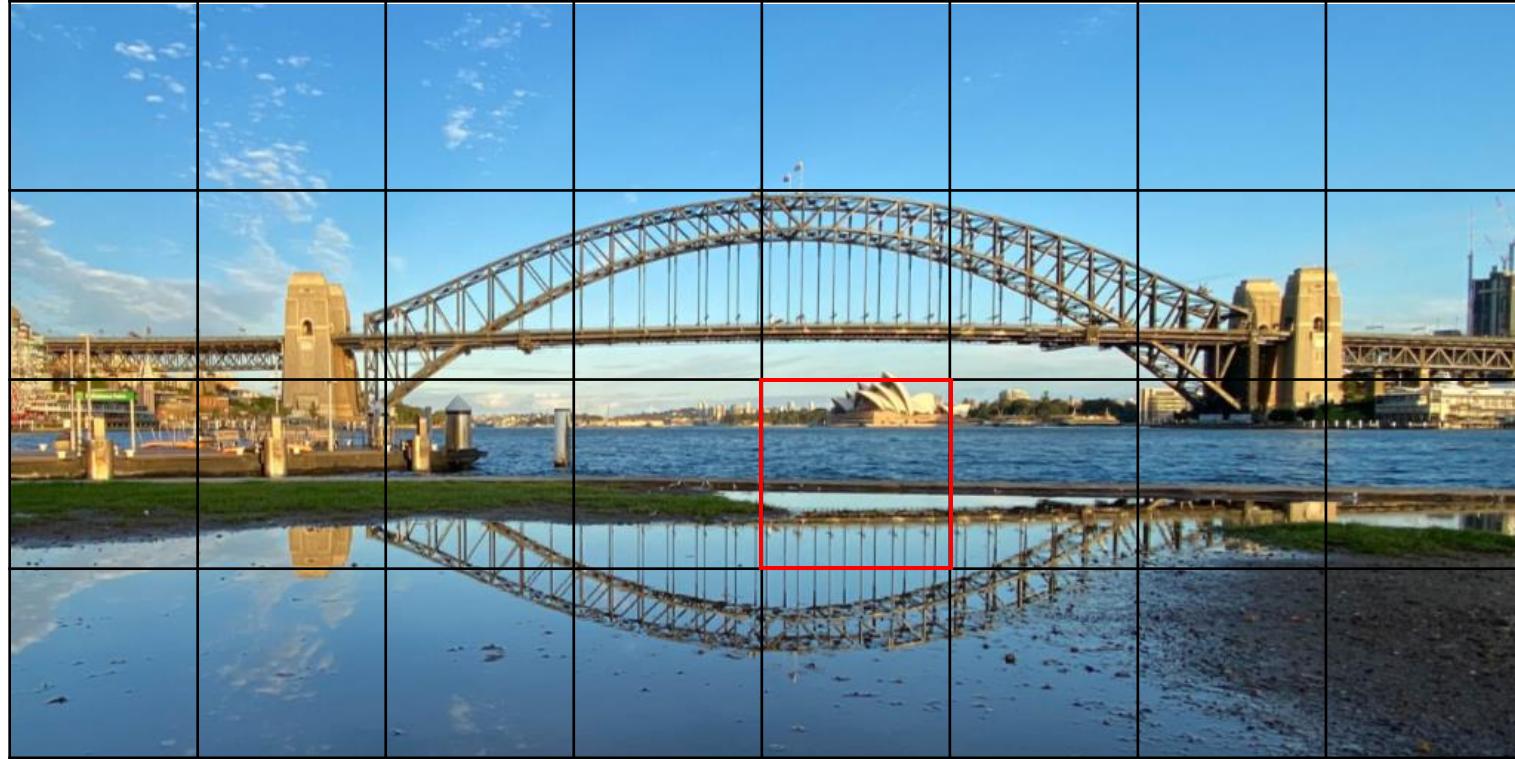


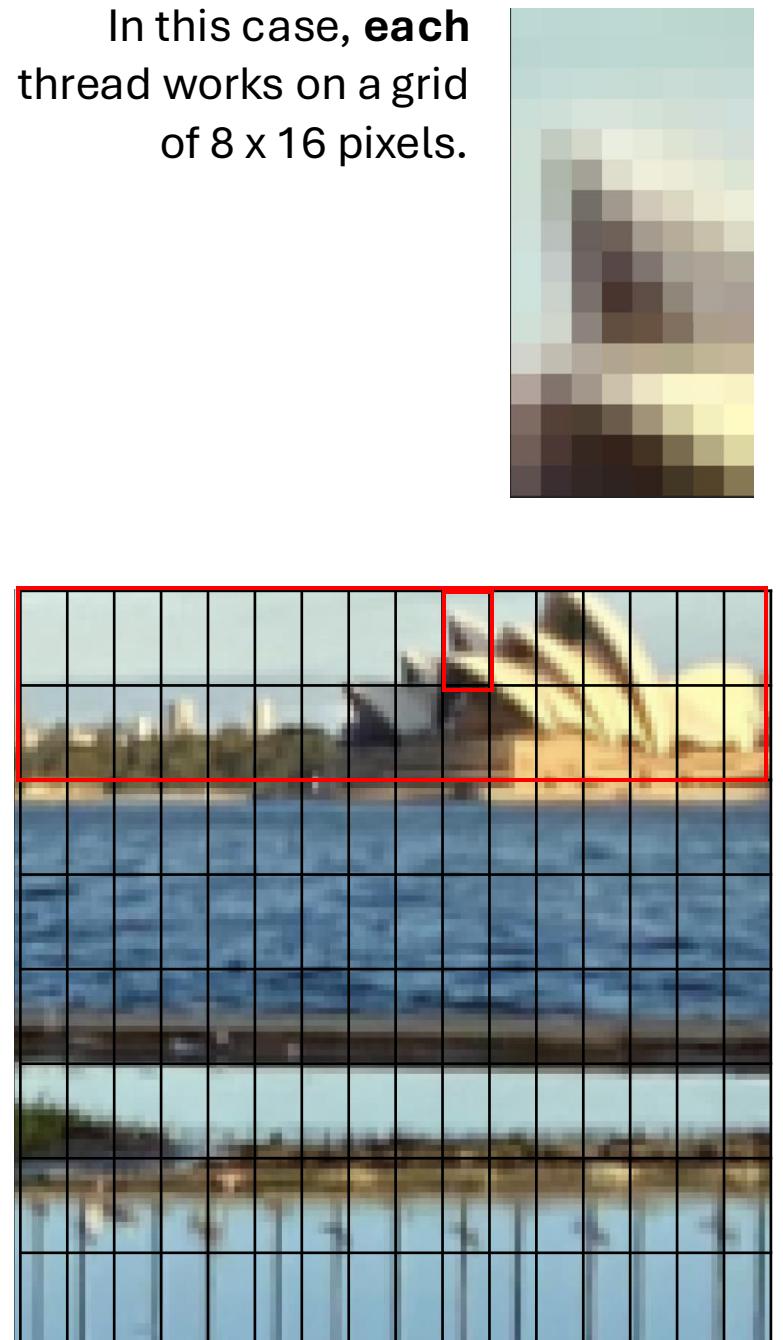
Image 1024 x 512 pixels

Split into 8 x 4 ‘threadgroups’,  
each of 16 x 8 threads, in this case.

OR...

Threads in a threadgroup  
execute in ‘warps’ of 32 threads  
running together in lockstep.

In this case, **each**  
thread works on a grid  
of 8 x 16 pixels.



```
float CalcMedian(float X[NPIXSQ_MAX], uint len)
{
    #define swap(a,b) {float t = X[a]; X[a] = X[b], X[b] = t;}
    uint left = 0, right = len - 1;
    uint cent = len/2;
    float pivot;
    uint pos, i;
    float median = 0.0;

    while (left < right) {
        pivot = X[cent];
        swap(cent, right);
        for (i = pos = left; i < right; i++) {
            if (X[i] < pivot) {
                swap(i, pos);
                pos++;
            }
        }
        swap(right, pos);
        if (pos == cent) break;
        if (pos < cent) left = pos + 1;
        else right = pos - 1;
    }
    median = X[cent];
```

The first published picture of the Mandelbrot set, by Robert W. Brooks and Peter Matelski in 1978

By Elphaba - Own work, Public Domain,  
<https://commons.wikimedia.org/w/index.php?curid=1627040>